

CSE-3211 LAB 1 Part 2

Shadman Wadith

Roll :- 27

20,June,2021

Question and Answers

Question 1: What is the vm system called that is configured for assignment 0?

Answer : DUMBVM

Location: kern/arch/mips/conf/conf.arch

Question 2. Which register number is used for the stack pointer (sp) in OS/161?

Answer : 29

Location: kern/arch/mips/include/kern/regdefs.h

Question 3. What bus/busses does OS/161 support?

Answer : LAMEbus

Location: kern/arch/sys161/include/bus.h

Question 4. What is the difference between splhigh and spl0?

Answer : splhigh() sets IPL to the highest value, disabling all interrupts. spl0() sets IPL to 0, enabling all interrupts.

Location: kern/include/spl.h

Question 5. Why do we use typedefs like u_int32_t instead of simply saying "int"?

Answer: Since we're in a 32-bit platform, size_t, ssize_t, and ptrdiff_t can correctly be either (unsigned) int or (unsigned) long. However, if we don't define it to the same one gcc is using, gcc will get upset. So, in order to make sure that we really get a 32-bit unsigned integer as unsigned int depends on the platform, we use u_int32_t.

Location: kern/arch/include/mips/kern/types.h

Question 6: What must be the first thing in the process control block?

Answer : Process state

Question 7. What does splx return?

Answer: Returns old spl level

Location: kern/thread/spl.c

Question 8. What is the highest interrupt level?

Answer: There are two interrupt levels, 0 and 1. 1 is defined as IPL HIGH.

Location: kern/include/spl.h

Question 9. What function is called when user-level code generates a fatal fault?

Answer: Function called when user-level code hits a fatal fault is static void kill curthread (vaddr t epc, unsigned code, vaddr tvaddr).

Location: kern/arch/mips/locore/trap.c

Question 10. How frequently are hardclock interrupts generated?

Answer: Hardclocks per second is 100Hz (#define HZ 100)

Location: kern/include/clock.h

Question 11. What functions comprise the standard interface to a VFS device?

Answer: devop_eachopen - called on each open call to allow denying the open. devop_io - for both reads and writes (the uio indicates the direction). devop_ioctl - miscellaneous control operations.

Location: kern/include/device.h

Question 12. How many characters are allowed in a volume name?

Answer: 32 (#define SFS VOLNAME SIZE 32 /*max length of volume name*/)

Location: kern/include/kern/sfs.h

Question 13. How many direct blocks does an SFS file have?

Answer: 15 (#define SFS NDIRECT 15 /*# of direct blocks in inode*/) Location:

kern/include/kern/sfs.h

Question 14. What is the standard interface to a file system (i.e., what functions must you implement to implement a new file system)?

Answer: fsop_sync - Flush all dirty buffers to disk. fsop_getvolname - Return volume name of filesystem. fsop_getroot - Return root vnode of filesystem. fsop_unmount - Attempt unmount of filesystem.

Location: kern/include/fs.h

Question 15. What function puts a thread to sleep?

Answer: void wchan_sleep(struct wchan *wc, struct spinlock *lk);

Location: kern/thread/thread.c

Question 16. How large are OS/161 pids?

Answer: 32 bit (typedef _i32_pid_t; /* Process ID */)

Location: kern/include/kern/types.h

Question 17. What operations can you do on a vnode?

Answer: `vop_eachopen`, `vop_reclaim`, `vop_read`, `vop_readlink` `vop_eachopen`, `vop_reclaim`, `vop_read`, `vop_readlink`, `vop_getdirentry`, `vop_write`, `vop_ioctl`, `vop_stat`, `vop_gettype`, `vop_isseekable`, `vop_fsync`, `vop_mmap`, `vop_truncate`, `vop_namefile`, `vop_creat`, `vop_symlink`, `vop_mkdir`, `vop_link`, `vop_remove`, `vop_rmdir`, `vop_rename`, `vop_lookup`, `vop_lookupparent`.

Location: `kern/include/vnode.h`

Question 18. What is the maximum path length in OS/161?

Answer : 1024

Location: `kern/include/kern/limits.h`

Question 19. What is the system call number for a reboot?

Answer : 119

Question 20. Where is `STDIN_FILENO` defined?

Answer: `#define STDIN_FILENO 0 /* Standard input */`

Location: `kern/include/kern/unistd.h`

Question 21. What does `kmain()` do?

Answer: Kernel main : Boot up, then fork the menu thread; wait for a reboot request, and then shut down.

Location: `kern/main/main.c`

Question 22. Is it OK to initialise the thread system before the scheduler? Why (not)?

Answer: Yes. Because the scheduler bootstrap creates the run queue and the thread bootstrap initializes the first thread.

Location: `kern/thread/thread.c`

Question 23. What is a zombie?

Answer: Zombies are threads that have exited but still need to have thread destroy called on them.

Location: `kern/thread/thread.c`

Question 24. How large is the initial run queue?

Answer: Though we did not find any initialization of the run queue, it seems that this line of code is an indication of run queue size (`curcpu - c_runqueue.tl_count = 0;`)

Location: `kern/thread/thread.c`

Question 25. What does a device name in OS/161 look like?

Answer: The name of a device is always just "device:". The VFS layer puts in the device name for us, so we don't need to do anything further. (e.g. "lhd0").

Location: `kern/vfs/vfslist.c`

Question 26. What does a raw device name in OS/161 look like?

Answer: Name of raw device (e.g., "lhd0raw"). Is non-NULL if and only if this device can have a filesystem mounted on it.

Location: kern/vfs/vfslist.c

Question 27. What lock protects the vnode reference count?

Answer: vn countlock

Location: kern/vfs/vnode.c

Question 28. What device types are currently supported?

Answer: Block Type and Character Type Devices.

Location: kern/vfs/device.c